

Problem statement:

I inherited a MongoDB database server with 60 collections and 100 or so indexes.

The business users are complaining about slow report completion times.

What can I do to improve performance ?

Scope:

System tuning-

- Memory
- Process
- Disk
- Network


Application tuning-

- Application architecture
- Statement design
- Data model design

Indexing, Query optimization

(Relational, so that we may compare/contrast)

- Detail command processing stages
- Can apply the 5 rules to a Rule Based query optimizer
- Apply 3 Index Negation guidelines
- Repair common query design problems-
 - Psuedo order by clause
 - OR topped queries
 - (And topped queries, non-compound)
- Drop and analyze query plans
- Articulate/control FJS, & ESR query processing patterns



What you will
leave with in 60
minutes:



Queries 1 & 2:

```
SELECT * FROM phonebook  
WHERE lastName LIKE "?son";
```

```
SELECT * FROM phonebook  
WHERE firstName = "Jennifer";
```

Query 3:

```
CREATE TABLE t1 (col1, col_2, .. 80 more columns);  
CREATE INDEX i1 ON t1 (col_gender);
```

```
SELECT * FROM t1  
WHERE col_gender = "F" AND col_age > 50;
```

Negation of an index

1. Non-initial substring
2. Non-anchored compound/composite key
3. (Poor) selectivity of a filter

Partial negation of an index-

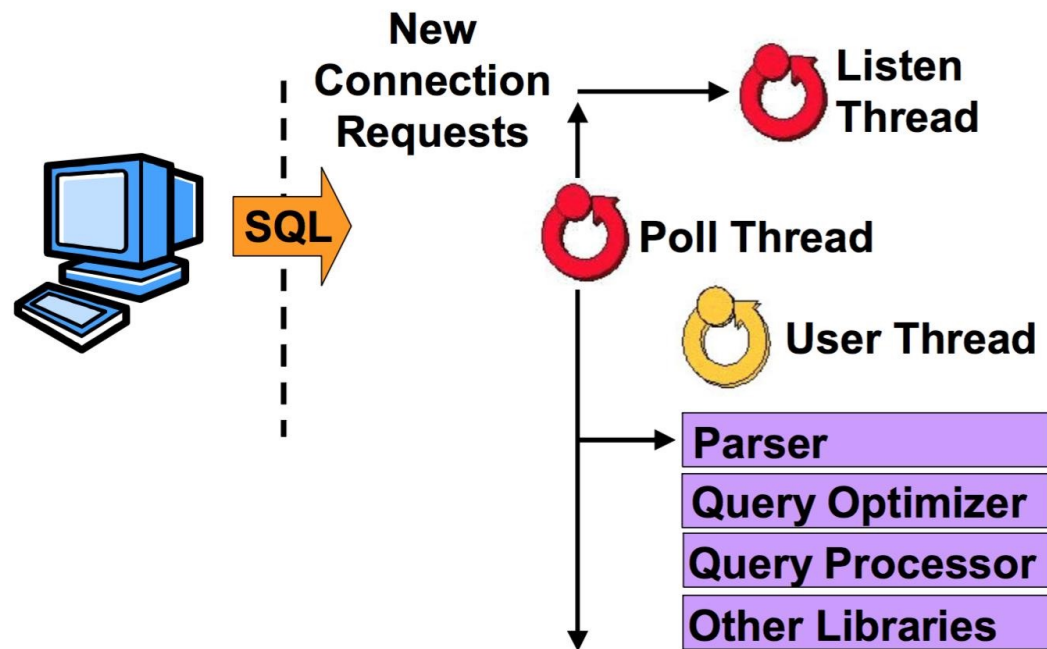
```
CREATE INDEX i1 ON t2 (col1, col2);  
//  
SELECT * FROM t1  
WHERE col1 > 100 AND col2 = "x";
```




Exception to all above-

Covered query/key-only

(n) Stage database server back end



Query Optimizers



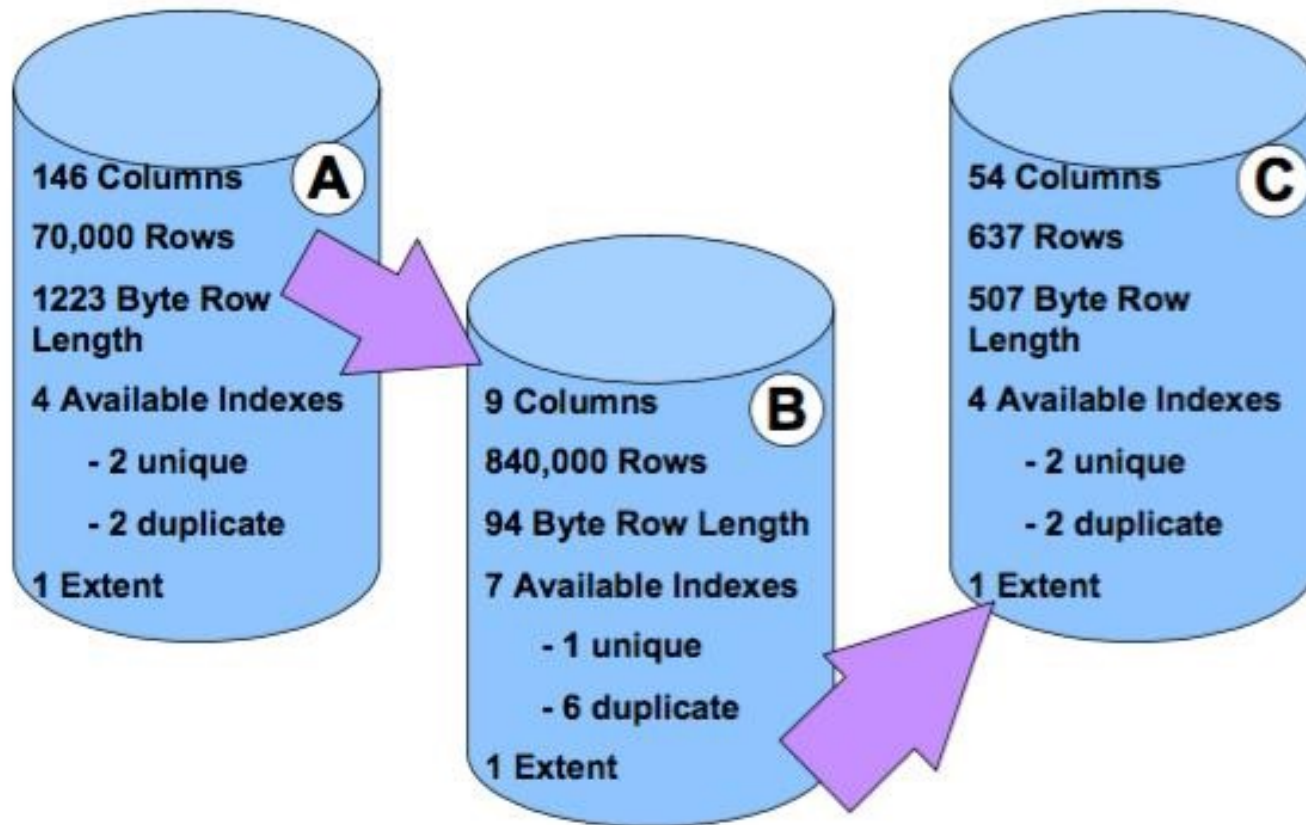
5 Rules to a rule based optimizer

1. Outer table joins
2. (Non-outer) table joins
3. Filter criteria (predicates)
4. Table size
5. Table cardinality

```
SELECT *  
FROM orderHeader, orderLineItems  
WHERE  
    oh.orderNumber =  
    oi.orderNumber;
```

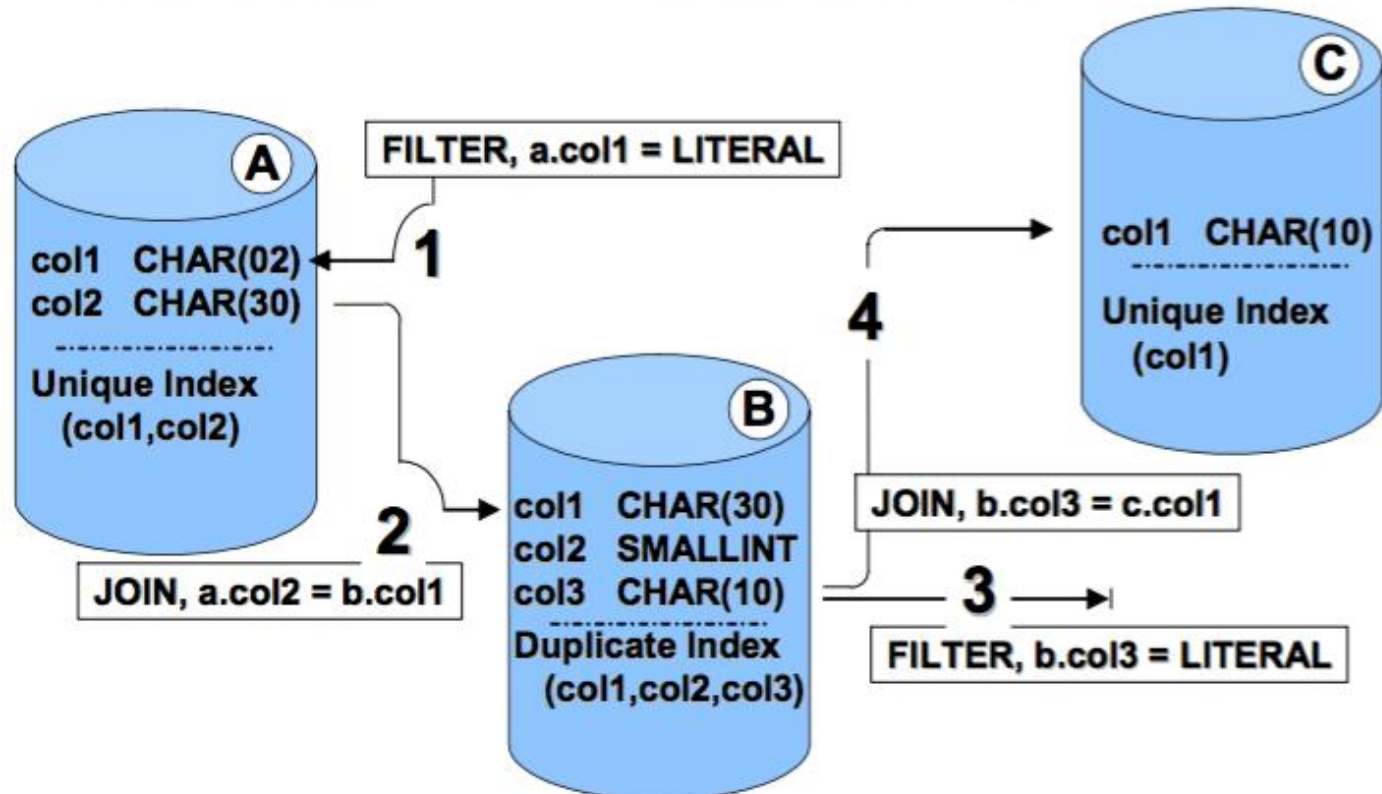
```
SELECT *  
FROM persons, OUTER automobiles  
WHERE  
    p.personId = a.personId;
```

Query 4: final/larger example using SQL

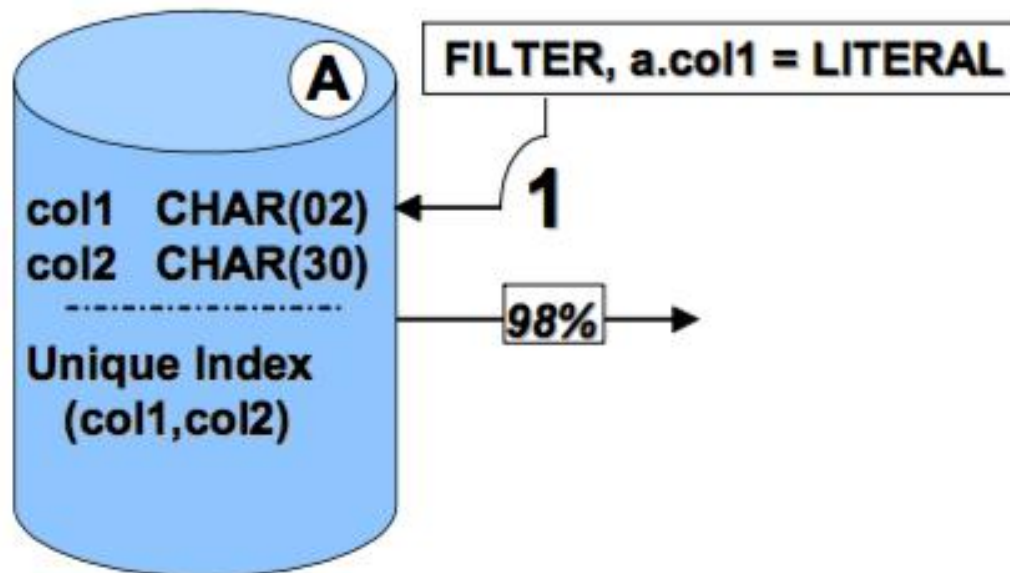


Example 4: query

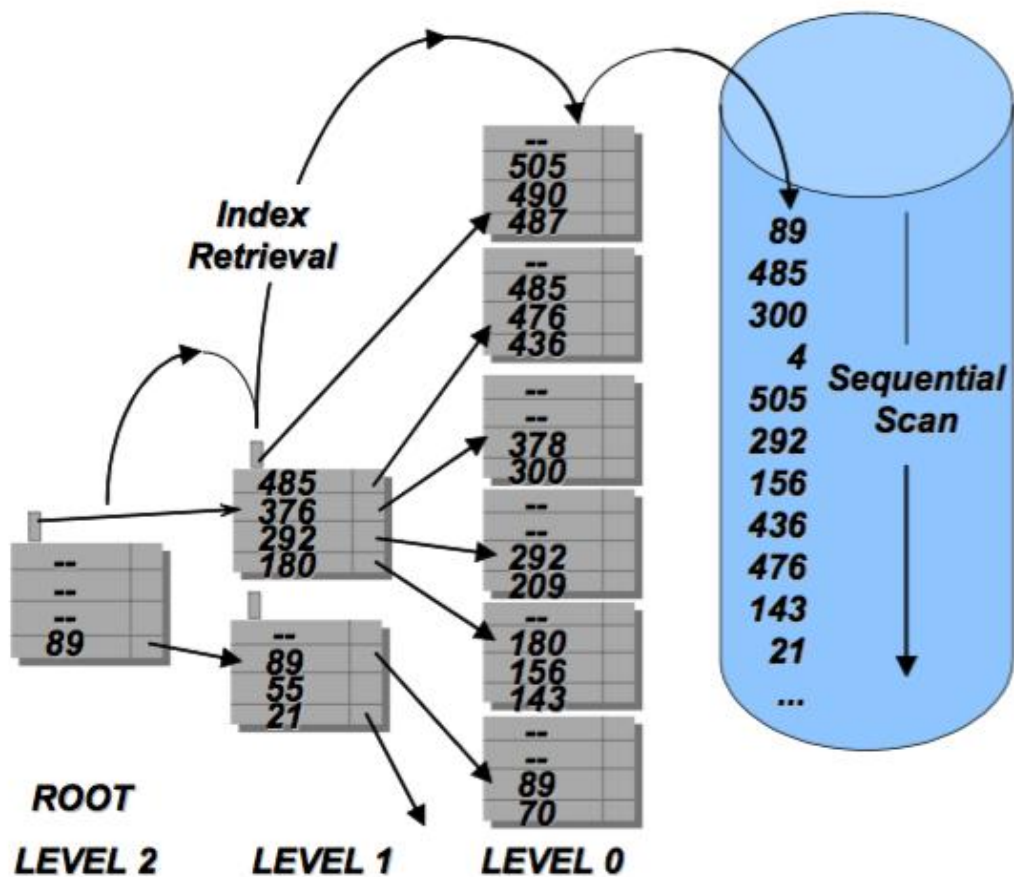
The query plan followed the table order, Tables "a", "b", then "c".



Query 4: First predicate




Collection access method: collection scan versus index scan



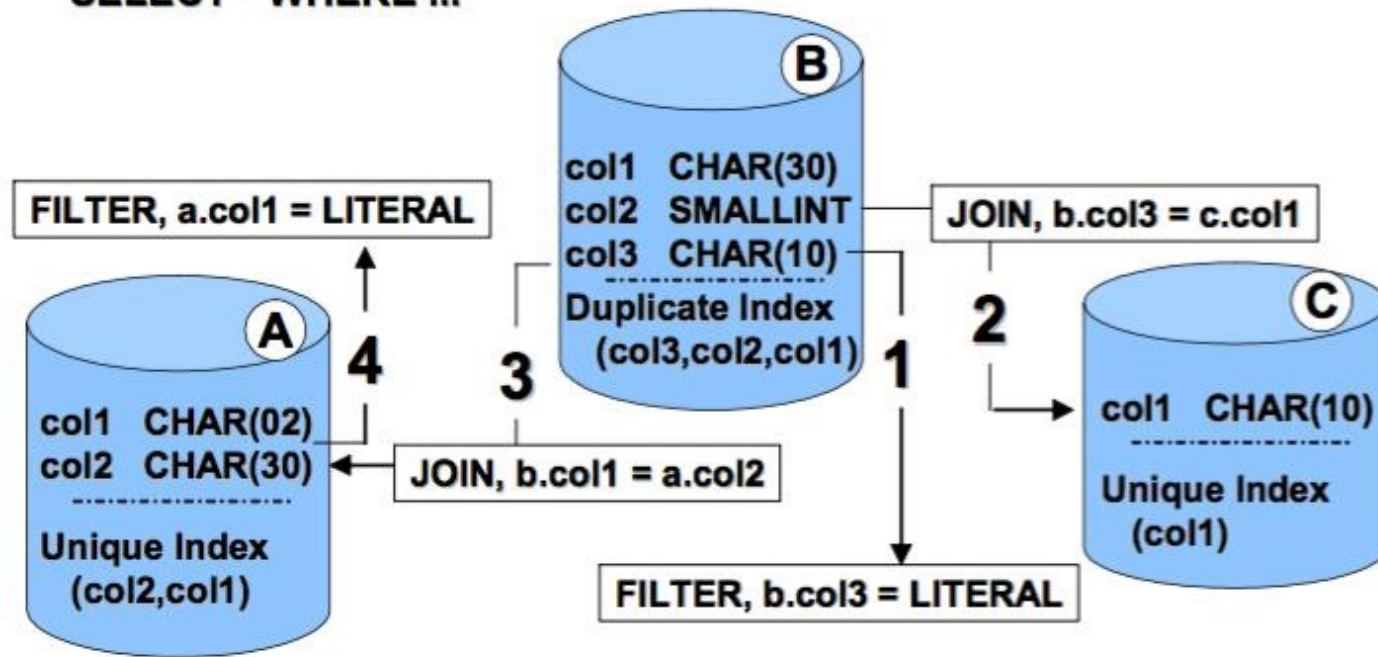
Query 4: Join and predicate

SELECT a.col1, COUNT(*) ... GROUP BY 1



Query 4: Optimal plan

The query plan followed the table order, Tables “b”, “c”, and then “a”.
SELECT * WHERE ...



FJS versus ESR: MongoDB

```
SELECT *  
FROM collection  
WHERE  
  col1 = 'x' and col2 > 'y'  
ORDER BY col3;
```

Filter -> Join -> Sort (FJS)

Equality -> Sort -> Range (ESR)

Problem statement:

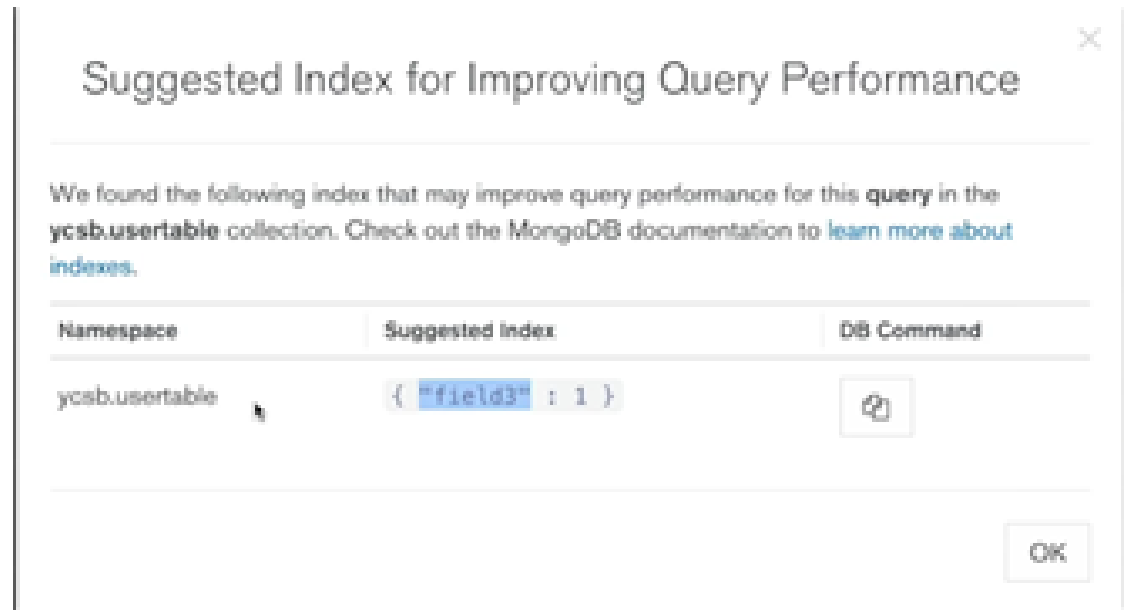
I inherited a MongoDB database server with 60 collections and 100 or so indexes.

The business users are complaining about slow report completion times.

What can I do to improve performance ?

Skills/tools we need:

- Which server
- Which logfile
- **Server profiling Level**
- Which queries
 - Cloud/Ops Manager !!
 - mtools
 - Text processing
- Drop the query plan
- Analyze the query plan
- Which indexes get used
- Other



Sample data set: zips.json

```
> db.zips.findOne()
{
  "_id" : ObjectId("570 .. 1c1f2"),
  "city" : "ACMAR",
  "zip" : "35004",
  "loc" : {
    "y" : 33.584132,
    "x" : 86.51557
  },
  "pop" : 6055,
  "state" : "AL"
}
> db.zips.count()
29470
```

```
>db.zips.find( { "state" : "WI", "pop" :
{ "$lt" : 50 } } ).sort( { "city" : 1 } )
```

Query 5: Dumping the query plan

```
db.zips.find( { "state" : "WI",  
  "pop" : { "$lt" : 50 } } ).sort(  
  { "city" :  
1 } ).explain("executionStats")
```

```
"winningPlan" : {  
  "stage" : "SORT", "sortPattern" : { "city" : 1 },  
  "inputStage" : {  
    "stage" : "COLLSCAN",  
    "filter" : {  
      "$and" : [ "state" : { "$eq" : "WI" "pop" : { "$lt" : 50  
"rejectedPlans" : [ ]  
"executionStats" : {  
  "nReturned" : 4,  
  "executionTimeMillis" : 16,  
  "totalKeysExamined" : 0,  
  "totalDocsExamined" : 29470,
```

Query 5: get indexes

```
> db.zips.getIndexes()  
[ {  
  "v" : 1, "key" : { "_id" : 1 },  
  "name" : "_id_",  
  "ns" : "test_db.zips"  
}]
```

```
db.zips.createIndex( { "state" : 1 , "pop" : 1 } )
```

Query 5: attempt 2

```
winningPlan" : {  
  "stage" : "SORT", "sortPattern" : { "city" : 1 },  
  "inputStage" : {  
    "stage" : "FETCH",  
    "inputStage" : {  
      "stage" : "IXSCAN",  
      "keyPattern" : { "state" : 1, "pop" : 1 },  
      "indexBounds" : {  
        "state" : [ "[\"WI\", \"WI\"]" ],  
        "pop" : [ "[-inf.0, 50.0)" ]  
      }  
    }  
  }  
  "executionStats" : {  
    "nReturned" : 4,  
    "executionTimeMillis" : 1,  
    "totalKeysExamined" : 4,  
    "totalDocsExamined" : 4,  
  }  
}
```

```
db.zips.createIndex( { "state" : 1 , "city" : 1 , "pop" : 1 } )
```

```
"winningPlan" : {  
  "stage" : "SORT", "sortPattern" : { "city" : 1  
  "inputStage" : {  
    "stage" : "FETCH",  
  "inputStage" : {  
    "stage" : "IXSCAN",  
      "keyPattern" : { "state" : 1, "pop" : 1  
"rejectedPlans" : [  
  "stage" : "FETCH",  
  "inputStage" : {  
    "stage" : "IXSCAN",  
      "keyPattern" : { "state" : 1, "city" : 1, "pop" : 1 },  
      "indexBounds" : {  
        "state" : [ "[\"WI\", \"WI\"]" ],  
        "city" : [ "[MinKey, MaxKey]" ],  
        "pop" : [ "[-inf.0, 50.0)" ]
```

Query 5: attempt 3

Query 6: (pseudo order by clause)



```
db.zips.find( { "state" : "CO" }  
              ).sort( { "pop" : 1 } )
```

```
SELECT * FROM t1  
WHERE col1 = 'x'  
ORDER BY col2;
```

```
SELECT * FROM t1  
WHERE col1 = 'x'  
ORDER BY col1, col2;
```

```
SELECT * FROM t1  
WHERE col1 = 'x'  
ORDER BY 'x', col2;
```


Query 6: query plan

```
"winningPlan" : {  
  "stage" : "FETCH",  
  "inputStage" : {  
    "stage" : "IXSCAN",  
    "keyPattern" : { "state" : 1, "pop" : 1 },  
    "indexBounds" : {  
      "state" : [ "[\"CO\\", \"CO\"]" ],  
      "pop" : [ "[MinKey, MaxKey]" ]  
    }  
  }  
  "rejectedPlans" : [  
    "stage" : "SORT",  
    "sortPattern" : { "pop" : 1 },  
    "inputStage" : {
```



```
    "stage" : "FETCH",  
    "inputStage" : {  
      "stage" : "IXSCAN",  
      "keyPattern" : { "state" : 1, "city" : 1, "pop" : 1 },  
      "indexBounds" : {  
        "state" : [ "[\"CO\\", \"CO\"]" ],  
        "city" : [ "[MinKey, MaxKey]" ],  
        "pop" : [ "[MinKey, MaxKey]" ]  
      }  
    }  
  }  
  "executionStats" : {  
    "nReturned" : 416,  
    "executionTimeMillis" : 1,  
    "totalKeysExamined" : 416,  
    "totalDocsExamined" : 416,
```

Review the indexes we have so far

```
> db.zips.getIndexes()
```

```
_id
```

```
db.zips.createIndex( { "state" : 1 , "pop" : 1 } )
```

```
db.zips.createIndex( { "state" : 1 , "city" : 1 , "pop" : 1 } )
```

Query 7: OR topped query

```
db.zips.find( { "$or" : [ { "state" : "UT" }, { "pop" : 2 } ] } )
```


```
"winningPlan" : {  
  "inputStage" : {  
    "stage" : "COLLSCAN",  
    "filter" : {  
      "$or" : [  
        "pop" : { "$eq" : 2 }  
        "state" : { "$eq" : "UT" }  
      ]  
    }  
  }  
  "rejectedPlans" : [ ]  
  "executionStats" : {  
    "nReturned" : 215,  
    "executionTimeMillis" : 22,  
    "totalKeysExamined" : 0,  
    "totalDocsExamined" : 29470,
```

```
SELECT * FROM t1  
WHERE  
  order_date = TODAY  
OR  
  ship_weight < 10;
```

Query 7: solution

```
db.zips.createIndex( { "pop" : 1 } )
```

```
"winningPlan" : {  
  "inputStage" : {  
    "stage" : "FETCH",  
    "inputStage" : {  
      "stage" : "OR",  
      "inputStages" : [  
        "stage" : "IXSCAN",  
        "keyPattern" : { "state" : 1, "pop" : 1 },  
        "indexBounds" : {  
          "state" : [ "[\"UT\\", \"UT\\"]" ],  
          "pop" : [ "[MinKey, MaxKey]" ]  
        }  
      ]  
    }  
  }  
}
```



```
"stage" : "IXSCAN",  
  "keyPattern" : { "pop" : 1 },  
  "indexBounds" : {  
    "pop" : [ "[2.0, 2.0]" ]  
  },  
  "rejectedPlans" : [ ]  
"executionStats" : {  
  "nReturned" : 215,  
  "executionTimeMillis" : 2,  
  "totalKeysExamined" : 215,  
  "totalDocsExamined" : 215,
```

Topics not previously covered



- How to tell which indexes are being used
 - How to tell if an index is unique
 - Smoke tests
 - Covered queries
 - MongoDB index types
 - When do winning query plans get evacuated
 - Index intersection
 - Building indexes (online/offline)
 - Sharding and queries, query plans
 - Capped collections, tailable cursors
 - Optimizer hints
 - Memory limits
 - Query rewrite (aggregation pipeline optimization)
-
- Which server
 - Which logfile
 - (Server profiling Level)
 - Which queries
 - mtools
 - Text processing

Resources:

- The parent to this preso,
<https://github.com/farrell0/MongoDB-Developers-Notebook>
- An excellent query primer (110 pages)
<http://www.redbooks.ibm.com/abstracts/sg247138.html?Open>
(Chapters 10 and 11.)
- University.MongoDB.com
- https://www.mongodb.com/consulting-test#performance_evaluation
- [zips.json](http://media.mongodb.org/zips.json)
<http://media.mongodb.org/zips.json>
- Call Dave Lutz, at home, On Sunday (early)
(512)555/1212

Backup Slides

How to tell which indexes are being used

```
db.zips.aggregate( [ { "$indexStats" : {} } ] ).pretty()
{ "name" : "pop_1",
  "key" : { "pop" : 1 },
  "host" : "rhhhost00.grid:27017",
  "accesses" : {
    "ops" : NumberLong(15),
    "since" : ISODate("2016-04-19T07:13:44.546Z") } }
{ "name" : "state_1_city_1_pop_1",
  "key" : { "state" : 1, "city" : 1, "pop" : 1 },
  "host" : "rhhhost00.grid:27017",
  "accesses" : {
    "ops" : NumberLong(0),
    "since" : ISODate("2016-04-19T06:49:11.765Z") } }
```


How to tell if an index is unique

```
db.t1.createIndex( { "k1" : 1 },
  { "unique" : true })
{
  "createdCollectionAutomatically" : true,
  "numIndexesBefore" : 1,
  "numIndexesAfter" : 2,
  "ok" : 1
}
```

```
> db.t1.getIndexes()
[ { "v" : 1,
  "key" : { "_id" : 1 },
  "name" : "_id_",
  "ns" : "test_db.t1" },
  { "v" : 1,
  "unique" : true,
  "key" : { "k1" : 1 },
  "name" : "k1_1",
  "ns" : "test_db.t1" }
]
```

Smoke tests

Every night, gather a set of statistics about your hard disk fullness, and about the performance of a set of queries that are strategic to the application.

For queries we wish to record-

- The number of documents returned
- The winning query plan
- Elapsed time, disk and memory consumed
- Other

```
> db.zips.find( { "pop" : { "$lt" : 200 } },
  { "_id" : 0, "pop" : 1 } ).sort(
  { "pop" : -1 } ).explain()
"winningPlan" : {
  "stage" : "PROJECTION",
  "transformBy" : {
    "_id" : 0,
    "pop" : 1 },
  "inputStage" : {
    "stage" : "IXSCAN",
    "keyPattern" : { "pop" : 1 },
    "indexBounds" : {
      "pop" : [ "(200.0, -inf.0]" ]
    }
  }
  "rejectedPlans" : [ ]
}
```

Covered queries

When does the winning plan get evacuated

In short, the cached query plan is re-evaluated if:

- The collection receives 1000 or more writes
- An index is added or dropped
- A reindex operation is performed
- mongod is restarted
- You run a query with explain

Index intersection

```
db.zips.find( { "$or" : [ { "state" : "UT" }, { "pop" : 2 } ] } )
```

```
db.zips.find( { "city" : "EAST TROY", "zip" : 53120 } )
```

Building indexes

```
db.zips.createIndex( { "zip" : 1 }, { "background" : true } )
```

Capped collections

```
db.createCollection("my_collection",  
{ capped : true, size : 5242880,  
max : 5000 } )
```

```
from pymongo import Connection  
import time
```

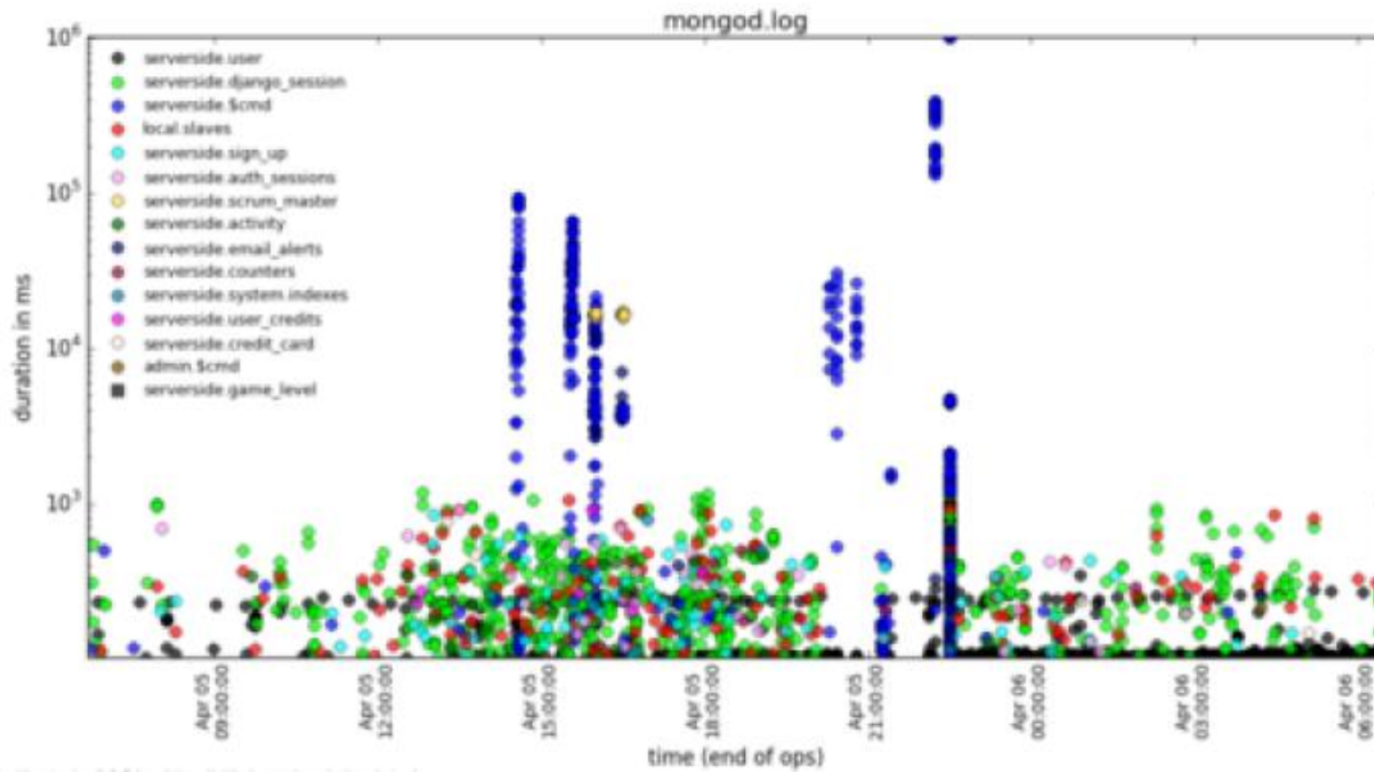
```
db = Connection().my_db  
coll = db.my_collection  
cursor = coll.find(tailable=True)  
while cursor.alive:  
    try:  
        doc = cursor.next()  
        print doc  
    except StopIteration:  
        time.sleep(1)
```

Memory limits: 32 MB, 100 MB

```
"executionStages" : {  
  "stage" : "SORT",  
  "nReturned" : 1,  
  "executionTimeMillisEstimate" : 60,  
  ...  
  "sortPattern" : { "City" : 1 },  
  "memUsage" : 120,  
  "memLimit" : 33554432,
```

```
db.zips.aggregate([  
  { "$group" :  
    { "_id" : "$state",  
      //  
      "totalPop" : { "$sum" : "$pop" },  
      "cityCount" : { "$sum" : 1 }  
    }  
  },  
  { "$sort" : { "_id" : 1 } }  
],  
  { "allowDiskUse" : true }  
)
```


mtools: mplotquery



But first:

Y/N I have more than 24 months
experience with SQL


Y/N I have more than 6 months
experience with MongoDB

Y/N I have dropped a MongoDB
explain plan, understood it,
made changes, and was happy

Y/N Puppies scare me



Two more examples: Queries 8 and 9

find()  **aggregate()**

- optimizer hints
- \$lookup()

Query 8: automatic query rewrite

```
> db.zips.aggregate(  
...  [  
...    { "$sort" :  
...      { "state" : 1 }  
...    },  
...    {  
...      "$match" :  
...        { "state" : { "$gt" : "M" } }  
...    }  
...  ],  
...  { "explain" : true } )
```

Query 8: Explain plan

```
"stages" : [
  "$cursor" : {
    "sort" : { "state" : 1 },
  },
  "winningPlan" : {
    "stage" : "FETCH",
    "inputStage" : {
      "stage" : "IXSCAN",
      "indexName" : "state_1_city_1",
      "indexBounds" : {
        "state" : [ ("M", {}) ],
        "city" : [ "[MinKey, MaxKey]" ]
      }
    }
  },
  "rejectedPlans" : [ ]
]
```

Query 9: Optimizer hints

```
> db.zips.find( { "city" : "EAST TROY" }).hint(
  { "zip" : 1 } ).explain("executionStats")
```

```
  "winningPlan" : {
    "stage" : "FETCH",
    "filter" : { "city" : { "$eq" : "EAST TROY" } },
    "inputStage" : {
      "stage" : "IXSCAN", "keyPattern" : { "zip" : 1 },
      "indexBounds" : { "zip" : [ "[MinKey, MaxKey]" ] }
    }
  },
  "rejectedPlans" : [ ]
"executionStats" : {
  "nReturned" : 1,
  "executionTimeMillis" : 28,
  "totalKeysExamined" : 29470,
  "totalDocsExamined" : 29470,
```